

CMPE 256 HomeWork 1

```
In [1]: ▶ import pandas as pd
        from collections import Counter
        from sklearn.metrics.pairwise import cosine_similarity
        from surprise import SVD
        import numpy as np
        import surprise
        from surprise import Reader, Dataset

        from scipy import sparse
```

```
In [2]: ▶ import os
        import pandas as pd
        import numpy as np
        from surprise import Reader
        from surprise import Dataset
        from surprise.model_selection import KFold
        from surprise.model_selection import cross_validate
        from surprise import NormalPredictor
        from surprise import KNNBasic
        from surprise import KNNWithMeans
        from surprise import KNNWithZScore
        from surprise import KNNBaseline
        from surprise import SVD
        from surprise import SVDpp
        from surprise import NMF
        from surprise import SlopeOne
        from surprise import CoClustering
        from surprise.accuracy import rmse
        from surprise import accuracy
        from surprise.model_selection import train_test_split
        from surprise.model_selection import GridSearchCV
        from collections import defaultdict
```

1. Loading Training Data

```
In [3]: ▶ df=pd.read_csv('train.dat', sep='\t', header=None)
        df.columns=['user', 'movie', 'rating', 'timestamp']
```

```
In [4]: df.head()
```

Out[4]:

	user	movie	rating	timestamp
0	905	470	1	889325071
1	697	1518	5	879835275
2	855	1687	5	875638677
3	950	1447	5	877420720
4	806	1170	4	879889337

```
In [5]: df.shape
```

Out[5]: (85724, 4)

2. Defining the rating scale and building the dataset

```
In [6]: reader = Reader(rating_scale=(0.5, 5.0))
data_set = Dataset.load_from_df(df[['user', 'movie', 'rating']], reader)
```

3. Trying with baseline algorithms

```
In [7]: #Reference from: https://towardsdatascience.com/movie-recommender-system-part-1
results = []

for alg in [SVD(), NMF(), NormalPredictor(), KNNBasic()]:

    res = cross_validate(alg, data_set, measures=['RMSE'], cv=3, verbose=False)
    tmp = pd.DataFrame.from_dict(res).mean(axis=0)
    tmp = tmp.append(pd.Series([str(alg).split(' ')[0].split('.')[1]], index=[0]))
    results.append(tmp)
```

```
Computing the msd similarity matrix...
Done computing similarity matrix.
Computing the msd similarity matrix...
Done computing similarity matrix.
Computing the msd similarity matrix...
Done computing similarity matrix.
```

```
In [8]: ▶ surprise_results = pd.DataFrame(results).set_index('Algorithm').sort_values('surprise_results')
```

Out[8]:

	test_rmse	fit_time	test_time
Algorithm			
SVD	0.950236	2.684577	0.159854
NMF	0.987491	3.169755	0.148271
KNNBasic	0.996197	0.207396	2.897233
NormalPredictor	1.522411	0.088283	0.211689

3. GridSearchCV to find best parameters in SVD

```
In [ ]: ▶ #Reference from: https://towardsdatascience.com/movie-recommender-system-part

param_grid = {'n_factors': [25, 30, 35, 40], 'n_epochs': [15, 20, 25], 'lr_all': [0.05, 0.1, 0.15],
              'reg_all': [0.08, 0.1, 0.15]}
gs = GridSearchCV(SVD, param_grid, measures=['rmse', 'mae'], cv=3)
gs.fit(data_set)
algo = gs.best_estimator['rmse']
print(gs.best_score['rmse'])
print(gs.best_params['rmse'])
```

```
In [9]: ▶ ## Previous best scores
## Score: 0.9301
##{'n_factors': 35, 'n_epochs': 30, 'lr_all': 0.008, 'reg_all': 0.05}
```

4. Loading test Data

```
In [10]: ▶ df_test=pd.read_csv('test.dat', sep='\t', header=None)
df_test.columns=['user', 'movie']
df_test['rating']=0
```

```
In [11]: df_test
```

```
Out[11]:
```

	user	movie	rating
0	158	951	0
1	521	1202	0
2	98	1556	0
3	292	1583	0
4	68	1064	0
...
2149	537	1414	0
2150	618	1448	0
2151	154	1519	0
2152	154	1429	0
2153	826	1602	0

2154 rows × 3 columns

```
In [12]: # Building test set
reader = Reader(rating_scale=(0.5, 5.0))
data_test = Dataset.load_from_df(df_test[['user', 'movie', 'rating']], reader)
```

5. Training on the whole training data

```
In [13]: train_data = data_set.build_full_trainset()

reg = SVD(n_factors=35, n_epochs=30, lr_all=0.008, reg_all=0.05)
reg.fit(train_data)
```

```
Out[13]: <surprise.prediction_algorithms.matrix_factorization.SVD at 0x1fc54038460>
```

```
In [14]: data_set.df.to_numpy()
```

```
Out[14]: array([[ 905,  470,   1],
                [ 697, 1518,   5],
                [ 855, 1687,   5],
                ...,
                [ 167, 1036,   3],
                [ 508, 1528,   3],
                [  76, 1586,   3]], dtype=int64)
```

```
In [15]: ▶ #RMSE for training data
preds_train = reg.test(data_set.df.to_numpy())
accuracy.rmse(preds_train)

RMSE: 0.6743
```

```
Out[15]: 0.6743489755060518
```

6. Pre-processing testing data

```
In [16]: ▶ testset = [data_test.df.loc[i].to_list() for i in range(len(data_test.df))]
```

7. Making predictions on pre-processed test.dat

```
In [17]: ▶ predictions = reg.test(testset)
```

```
In [18]: ▶ pred=[]
for prediction in predictions:
    pred.append(prediction[3])
```

```
In [19]: ▶ len(pred)
```

```
Out[19]: 2154
```

```
In [20]: ▶ submission_df=pd.DataFrame(pred)
```

```
In [22]: ▶ submission_df.to_csv('submission2.dat', index=False, header=False)
```

```
In [23]: ▶ submission=pd.read_csv('submission2.dat')
```

In [24]:  submission

Out[24]:

2.9401309514158003	
0	3.958190
1	3.172791
2	4.085288
3	3.391724
4	4.451186
...	...
2148	3.874339
2149	2.787361
2150	4.150590
2151	3.701678
2152	4.998140

2153 rows × 1 columns